

WATER POTABILITY PREDICTION

Step-1: Import all the required libraries which are used to train the model or visualise the data

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.impute import SimpleImputer
import warnings
warnings.filterwarnings("ignore")
import plotly
import plotly.offline as pyo
import plotly.express as px
import plotly.graph_objs as go
pyo.init_notebook_mode()
import plotly.figure_factory as ff
```

Step-2: Read the data

```
In [2]: df=pd.read_csv("/home/silpa/Downloads/water_potability.csv")
print(df.head())
```

| | ph | Hardness | Solids | Chloramines | Sulfate | Conductivity | \ |
|---|----------|------------|--------------|-------------|------------|--------------|---|
| 0 | NaN | 204.890455 | 20791.318981 | 7.300212 | 368.516441 | 564.308654 | |
| 1 | 3.716080 | 129.422921 | 18630.057858 | 6.635246 | NaN | 592.885359 | |
| 2 | 8.099124 | 224.236259 | 19909.541732 | 9.275884 | NaN | 418.606213 | |
| 3 | 8.316766 | 214.373394 | 22018.417441 | 8.059332 | 356.886136 | 363.266516 | |
| 4 | 9.092223 | 181.101509 | 17978.986339 | 6.546600 | 310.135738 | 398.410813 | |

| | Organic_carbon | Trihalomethanes | Turbidity | Potability |
|---|----------------|-----------------|-----------|------------|
| 0 | 10.379783 | 86.990970 | 2.963135 | 0 |
| 1 | 15.180013 | 56.329076 | 4.500656 | 0 |
| 2 | 16.868637 | 66.420093 | 3.055934 | 0 |
| 3 | 18.436524 | 100.341674 | 4.628771 | 0 |
| 4 | 11.558279 | 31.997993 | 4.075075 | 0 |

Step-3: To know number of rows and columns in the data set

```
In [3]: df.shape
```

```
Out[3]: (3276, 10)
```

Step-4: Checking for missing values

```
In [4]: df.isna().sum()
```

```
Out[4]: ph                491
Hardness                0
Solids                  0
Chloramines             0
Sulfate                 781
Conductivity            0
Organic_carbon          0
Trihalomethanes        162
Turbidity               0
Potability              0
dtype: int64
```

Step-5: To get some information about the data set

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3276 entries, 0 to 3275
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   ph                    2785 non-null   float64
1   Hardness              3276 non-null   float64
2   Solids                3276 non-null   float64
3   Chloramines           3276 non-null   float64
4   Sulfate               2495 non-null   float64
5   Conductivity          3276 non-null   float64
6   Organic_carbon        3276 non-null   float64
7   Trihalomethanes       3114 non-null   float64
8   Turbidity             3276 non-null   float64
9   Potability            3276 non-null   int64  
dtypes: float64(9), int64(1)
memory usage: 256.1 KB
```

Step-6: Describe the dataset which shows the minimum value, maximum value, mean value, count, standard deviation, etc.

```
In [6]: df.describe()
```

```
Out[6]:
```

| | ph | Hardness | Solids | Chloramines | Sulfate | Conductivity | Organic_carbon | Triha |
|-------|-------------|-------------|--------------|-------------|-------------|--------------|----------------|-------------|
| count | 2785.000000 | 3276.000000 | 3276.000000 | 3276.000000 | 2495.000000 | 3276.000000 | 3276.000000 | 3276.000000 |
| mean | 7.080795 | 196.369496 | 22014.092526 | 7.122277 | 333.775777 | 426.205111 | 14.284970 | 14.284970 |
| std | 1.594320 | 32.879761 | 8768.570828 | 1.583085 | 41.416840 | 80.824064 | 3.308162 | 3.308162 |
| min | 0.000000 | 47.432000 | 320.942611 | 0.352000 | 129.000000 | 181.483754 | 2.200000 | 2.200000 |
| 25% | 6.093092 | 176.850538 | 15666.690297 | 6.127421 | 307.699498 | 365.734414 | 12.065801 | 12.065801 |
| 50% | 7.036752 | 196.967627 | 20927.833607 | 7.130299 | 333.073546 | 421.884968 | 14.218338 | 14.218338 |
| 75% | 8.062066 | 216.667456 | 27332.762127 | 8.114887 | 359.950170 | 481.792304 | 16.557652 | 16.557652 |
| max | 14.000000 | 323.124000 | 61227.196008 | 13.127000 | 481.030642 | 753.342620 | 28.300000 | 28.300000 |

Step-7: Filling the missing values using a mean value of each feature.

```
In [7]: imputer = SimpleImputer(missing_values=np.nan, strategy='mean')
df2 = pd.DataFrame(imputer.fit_transform(df), columns=df.columns)
df2.isnull().sum()
```

```
Out[7]: ph                    0
Hardness                  0
Solids                    0
Chloramines               0
Sulfate                   0
Conductivity              0
Organic_carbon            0
Trihalomethanes           0
Turbidity                 0
Potability                0
dtype: int64
```


Step-8: Checking the value counts of our target feature Potability.

```
In [8]: df2.Potability.value_counts()
```

```
Out[8]: 0.0    1998  
        1.0    1278  
        Name: Potability, dtype: int64
```


Step-9: Visualising the potability using a countplot function.

```
In [9]: plt.figure(figsize=(7,7))  
sns.countplot(df2['Potability'])  
plt.xlabel('Potabiliy', fontsize=15)  
plt.ylabel('Count', fontsize=15)  
plt.show()
```



Step-10: visualising the pH value using a distplot function.

```
In [10]: plt.figure(figsize=(8,8))  
sns.distplot(df2['ph'], color='green')  
plt.xlabel('Ph', fontsize=15)  
plt.ylabel('Density', fontsize=15)  
plt.grid()  
plt.show()
```



```
In [11]: df2.columns
```


```
Out[11]: Index(['ph', 'Hardness', 'Solids', 'Chloramines', 'Sulfate', 'Conductivity',
               'Organic_carbon', 'Trihalomethanes', 'Turbidity', 'Potability'],
              dtype='object')
```

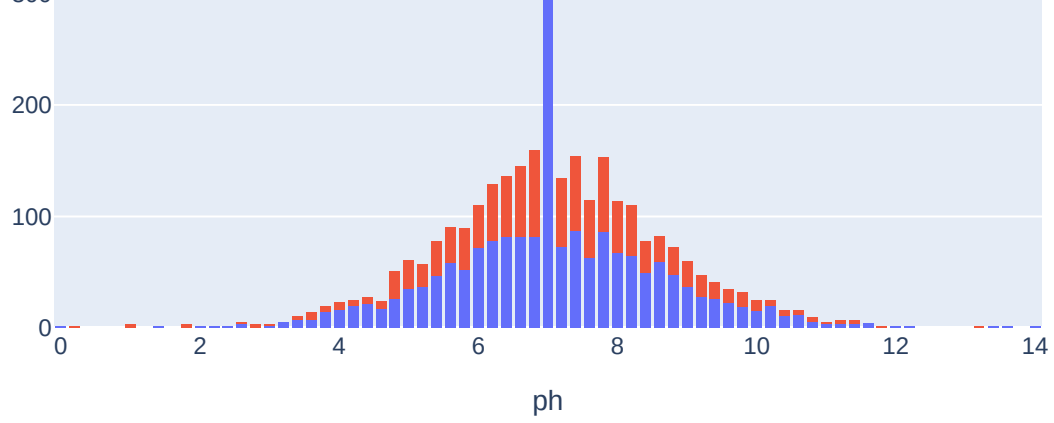
Distribution checking

```
In [12]: for colm in df2.columns:
          fig = px.histogram(df2, x=colm, color="Potability")
          fig.update_layout(title_text='The ' + colm + ' values of water according to the potabi
                              xaxis_title_text=colm,
                              yaxis_title_text='Count',
                              bargap=0.2,
                              bargroupgap=0.1)


          fig.show()
```

The ph values of water according to the potability category.




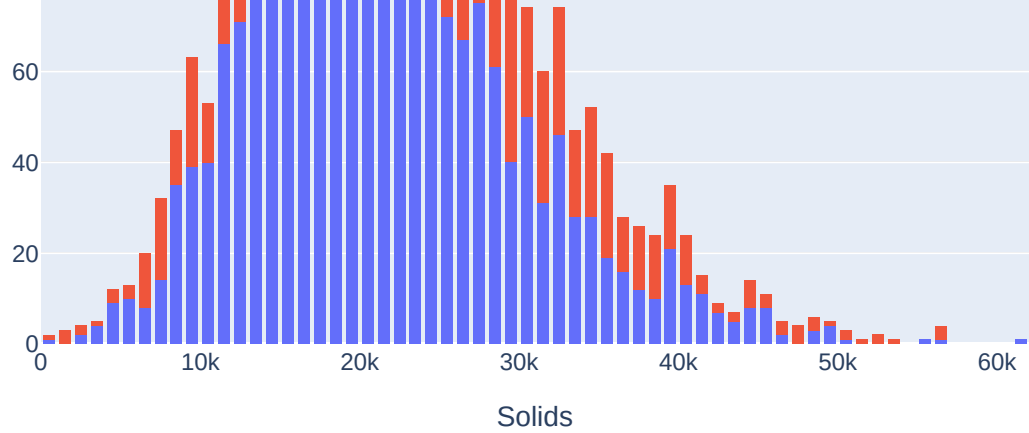


The Hardness values of water according to the potability category.




The Solids values of water according to the potability category.



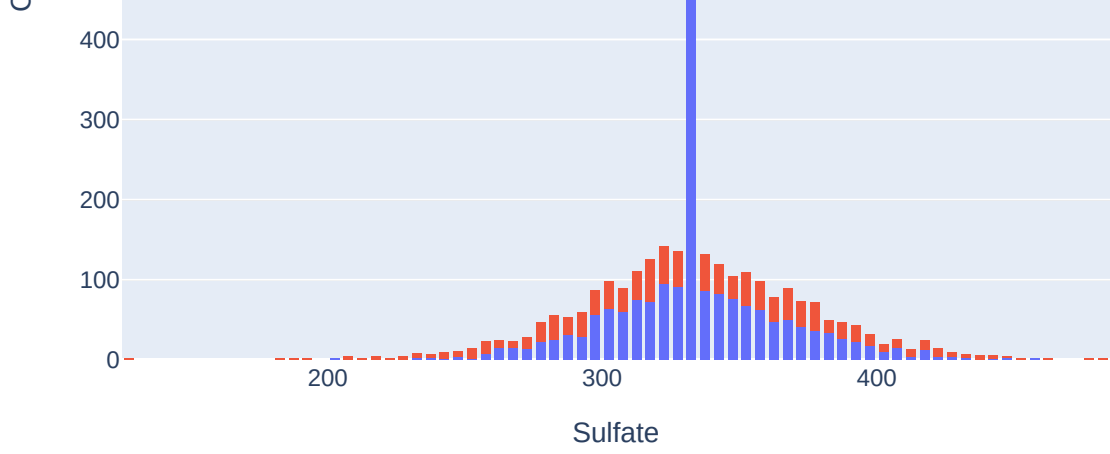


The Chloramines values of water according to the potability category.




The Sulfate values of water according to the potability category.




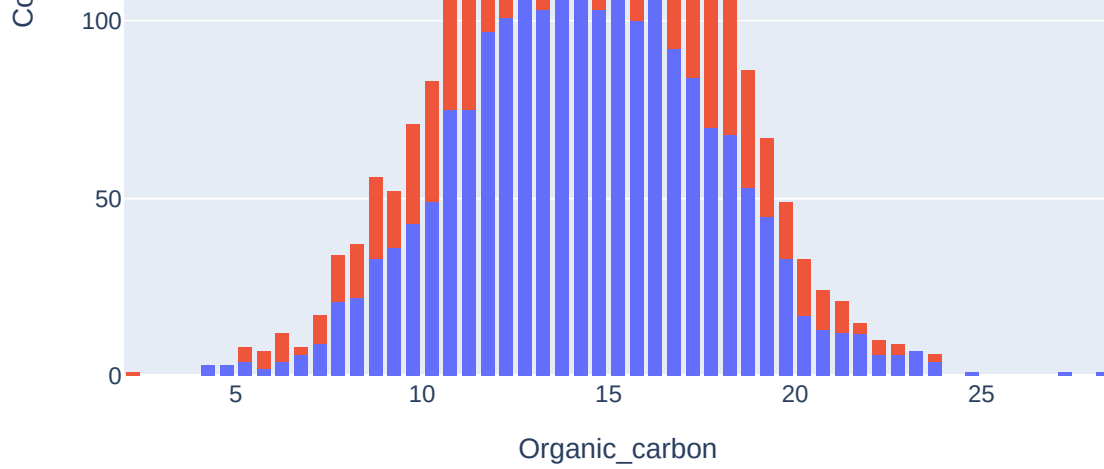


The Conductivity values of water according to the potability category.




The Organic_carbon values of water according to the potability category.







The Trihalomethanes values of water according to the potability category.




The Turbidity values of water according to the potability category.





The Potability values of water according to the potability category.



step-12: Assigning input and output

```
In [13]: x=df2.iloc[:, :-1].values
print(x)

[[7.08079450e+00  2.04890455e+02  2.07913190e+04  ...  1.03797831e+01
  8.69909705e+01  2.96313538e+00]
 [3.71608008e+00  1.29422921e+02  1.86300579e+04  ...  1.51800131e+01
  5.63290763e+01  4.50065627e+00]
 [8.09912419e+00  2.24236259e+02  1.99095417e+04  ...  1.68686369e+01
  6.64200925e+01  3.05593375e+00]
 ...
 [9.41951032e+00  1.75762646e+02  3.31555782e+04  ...  1.10390697e+01
```

```
6.98454003e+01 3.29887550e+00]
[5.12676292e+00 2.30603758e+02 1.19838694e+04 ... 1.11689462e+01
 7.74882131e+01 4.70865847e+00]
[7.87467136e+00 1.95102299e+02 1.74041771e+04 ... 1.61403676e+01
 7.86984463e+01 2.30914906e+00]]
```

```
In [14]: y=df2.iloc[:, -1].values
```

Step-13: Modelling

```
In [15]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30,random_state=101,strat
```

Normalization

```
In [16]: from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
scaler.fit(x_train)
x_train=scaler.transform(x_train)
x_test=scaler.transform(x_test)
```

KNN Algorithm

```
In [17]: from sklearn.neighbors import KNeighborsClassifier
cl=KNeighborsClassifier(n_neighbors=5)
cl.fit(x_train,y_train)
```

```
Out[17]: KNeighborsClassifier()
```

```
In [18]: y_pred=cl.predict(x_test)
print(cl.predict([[7.8,205.56,30700.55,5.2,400.56,66600.2,9.4,89.6,4.6]]))

[0.]
```

```
In [19]: from sklearn.metrics import classification_report,confusion_matrix,accuracy_score,plot_c
result=confusion_matrix(y_test,y_pred)
accuracy=accuracy_score(y_test,y_pred)
print(accuracy)
print(result)
r=classification_report(y_test,y_pred)
```

```
0.6408952187182095
[[472 128]
 [225 158]]
```

Naive Bayes Algorithm

```
In [20]: from sklearn.naive_bayes import GaussianNB
model=GaussianNB()
model.fit(x_train,y_train)
yg_pred=model.predict(x_test)
print(model.predict([[7.8,205.56,30700.55,5.2,400.56,66600.2,9.4,89.6,4.6]]))

[1.]
```

```
In [21]: cn=confusion_matrix(y_test,yg_pred)
print(cn)
acc=accuracy_score(y_test,yg_pred)
print(acc)
rpt=classification_report(y_test,yg_pred)

[[534 66]
 [302 81]]
```

Support Vector Machine Algorithm

```
In [22]: from sklearn.svm import SVC
classifier=SVC()
classifier.fit(x_train,y_train)
ys_pred=classifier.predict(x_test)
print(classifier.predict([[7.8,205.56,30700.55,5.2,400.56,66600.2,9.4,89.6,4.6]]))


[1.]
```

```
In [23]: cnf=confusion_matrix(y_test,ys_pred)
print(cnf)
ac=accuracy_score(y_test,ys_pred)
print(ac)
rprt=classification_report(y_test,ys_pred)

[[566  34]
 [277 106]]
0.6836215666327569
```

Comparison

```
In [24]: cm=confusion_matrix(y_test,ys_pred)
sns.heatmap(cm,annot=True,fmt=".0f")
plt.xlabel('y_pred')
plt.ylabel('y_test')
plt.title('Accuracy Score: {0}'.format(ac), size=12)
plt.show()
print('KNN:',r,"*" * 100,"\n","\n",'NAIVE_BAYES:',rpt,"*" * 100,"\n","\n",'SVC:',rprt)
```



| | | | | | |
|------|--------------|-----------|--------|----------|---------|
| KNN: | | precision | recall | f1-score | support |
| | 0.0 | 0.68 | 0.79 | 0.73 | 600 |
| | 1.0 | 0.55 | 0.41 | 0.47 | 383 |
| | accuracy | | | 0.64 | 983 |
| | macro avg | 0.61 | 0.60 | 0.60 | 983 |
| | weighted avg | 0.63 | 0.64 | 0.63 | 983 |

| | | | | | |
|--------------|-----|-----------|--------|----------|---------|
| NAIVE_BAYES: | | precision | recall | f1-score | support |
| | 0.0 | 0.64 | 0.89 | 0.74 | 600 |
| | 1.0 | 0.55 | 0.21 | 0.31 | 383 |

| | | | | |
|--------------|------|------|------|-----|
| accuracy | | | 0.63 | 983 |
| macro avg | 0.59 | 0.55 | 0.52 | 983 |
| weighted avg | 0.60 | 0.63 | 0.57 | 983 |

| | | | | | |
|--------------|------|-----------|--------|----------|---------|
| SVC: | | precision | recall | f1-score | support |
| | 0.0 | 0.67 | 0.94 | 0.78 | 600 |
| | 1.0 | 0.76 | 0.28 | 0.41 | 383 |
| accuracy | | | | 0.68 | 983 |
| macro avg | 0.71 | 0.61 | 0.59 | 983 | |
| weighted avg | 0.70 | 0.68 | 0.64 | 983 | |